

# deal.II and GPUs

Daniel Arndt  
Bruno Turcksin  
August 14, 2024

11th deal.II Users and Developers Workshop

ORNL is managed by UT-Battelle LLC for the US Department of Energy

## deal.II and GPUs

Places in deal.II that (can) use GPUs

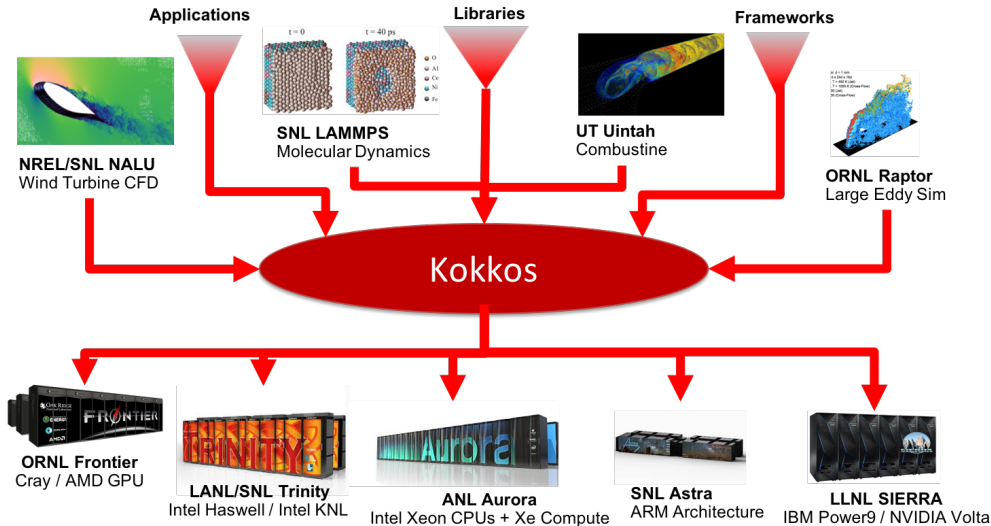
- PETSc (TBD)
- Trilinos
- `Portable::MatrixFree`
- `LinearAlgebra::distributed::Vector`
- `CUDAWrappers` (to be removed in the next release)

Most of these (except the latter two) use Kokkos under the hood!

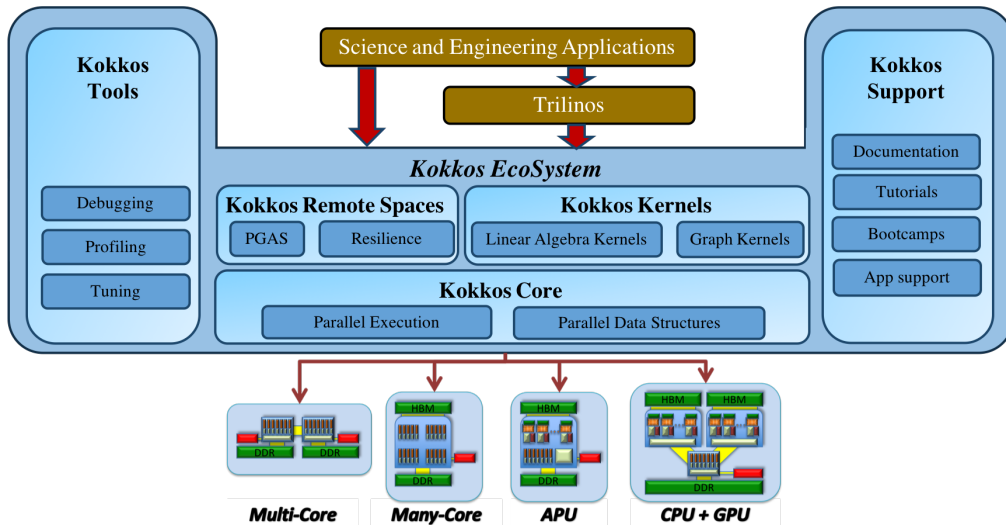
## What is Kokkos?

- A C++ Programming Model for Performance Portability
  - Template library on top CUDA, HIP, OpenMP, SYCL, ...
  - Aligns with developments in the C++ standard, e.g., `mdspan`, `atomic_ref`
- Expanding solution for common needs of modern science and engineering codes
  - Math libraries based on Kokkos
  - Tools for debugging, profiling and tuning
  - Interoperability with Fortran and Python
- Open Source project with a growing community
  - Maintained and developed at <https://github.com/kokkos>
  - Hundreds of users at many large institutions

# Kokkos as Portability Layer



# Kokkos EcoSystem





### Kokkos' basic capabilities:

- Simple 1D data parallel computational patterns
- Deciding where code is run and where data is placed
- Managing data access patterns for performance portability
- Multidimensional data parallelism

### Kokkos' advanced capabilities:

- Thread safety, thread scalability, and atomic operations
- Hierarchical patterns for maximizing parallelism
- Task-based programming with Kokkos

### Kokkos' tools and Kernels:

- How to profile, tune and debug Kokkos code
- Interacting with Python and Fortran
- Using KokkosKernels math library

# Kokkos Resources

## Online Resources:

- <https://github.com/kokkos:>
  - Primary Kokkos GitHub Organization
- <https://github.com/kokkos/kokkos-tutorials/wiki/Kokkos-Lecture-Series:>
  - Slides, recording and Q&A for the Lectures
- <https://kokkos.github.io/kokkos-core-wiki:>
  - Kokkos Core Wiki with API documentation
- <https://kokkosteam.slack.com:>
  - Slack channel for Kokkos.
  - Please join: fastest way to get your questions answered.
  - Can whitelist domains, or invite individual people.



# Kokkos iota

```
#include <Kokkos_Core.hpp>
int main() {
    Kokkos::ScopeGuard scope_guard;
    Kokkos::View<int*> view("view", 100);
    Kokkos::parallel_for(100, KOKKOS_LAMBDA(int i){view(i) = i;});
}
```

## Kokkos::View

- managed, multi-dimensional, reference-counted data container, or
- unmanaged, multi-dimensional accessor for external data
- motivated C++23 `std::mdspan`

## parallel\_reduce

```
double result;  
Kokkos::parallel_reduce(  
    Kokkos::RangePolicy(execution_space, start, end)),  
    KOKKOS_LAMBDA(int i, double& partial_sum) {  
        partial_sum += i;  
    }, result);
```

### Features:

- simple reductions (sum)
- multiple reductions per parallel construct
- custom reductions with arbitrary value types and reduction operations
- runtime sized array reductions
- pre- and post-callbacks for reductions (init, final)

## parallel\_scan

```
Kokkos::parallel_scan(  
    Kokkos::RangePolicy(execution_space, start, end),  
    KOKKOS_LAMBDA (const int index, value_type& update,  
                   const bool is_final) {  
        const value_type local_value = in_data(i);  
        // exclusive scan  
        if (is_final)  
            out_data_exclusive(i) = update;  
        update += local_value;  
        // inclusive scan  
        if (is_final)  
            out_data_inclusive(i) = update;  
    });
```

## MDRangePolicy

MDRangePolicy allows using up to a 6-dimensional index space with tiling

```
struct Functor{  
    KOKKOS_FUNCTION void operator(  
        int i, int j, int k, int l, int m) const { /*...*/ }  
};  
Kokkos::parallel_for(  
    Kokkos::MDRangePolicy(execution_space,  
        {s0,s1,s2,s3,s4}, {e0,e1,e2,e3,e4}),  
    Functor{});
```

## TeamPolicy

```
parallel_for("Label", TeamPolicy<>(numberOfTeams, teamSize, vectorLength),
    KOKKOS_LAMBDA (const member_type & teamMember) {
    /* beginning of outer body */
    parallel_for(TeamThreadRange(teamMember, thisTeamsRangeSize),
        [=] (const int indexWithinBatch[, ...]) {
        /* begin middle body */
        parallel_for(ThreadVectorRange(teamMember, thisVectorRangeSize),
            [=] (const int indexVectorRange) { /* inner body */});
        /* end middle body */
        });
    parallel_for(TeamVectorRange(teamMember, someSize),
        [=] (const int indexTeamVector) { /* nested body */});
    /* end of outer body */
});
```

## deal.II and GPUs

Places in deal.II that (can) use GPUs

- PETSc (TBD)
- Trilinos
- `Portable::MatrixFree`
- `LinearAlgebra::distributed::Vector`
- `CUDAWrappers` (to be removed in the next release)

All of these (except the latter) use Kokkos under the hood!

## Trilinos - Tpetra

Added support for Tpetra in this release (Jan-Philip Thiele and Sebastian Kinnewig)

- `Trilinos::TpetraWrappers::Vector<Number, MemorySpace>`
- `Trilinos::TpetraWrappers::SparseMatrix<Number, MemorySpace>`
- `Trilinos::TpetraWrappers::Solver*<Number, MemorySpace>`

Uses DualView: Memory is duplicated for `MemorySpace::Default`. All operations happen under the hood.

## Standalone linear algebra

`LinearAlgebra::distributed::Vector`

- `LinearAlgebra::distributed::Vector<Number, MemorySpace>`

Memory allocation only for the requested memory space, `MemorySpace::Host` (default) and `MemorySpace::Default`. Requires moving memory to and from the device explicitly.



# Assembly in deal.II

## Assembly approaches in deal.II

- Manually write the assembly of the matrix
- `MatrixCreator`
- `MeshWorker/MeshLoop`
- `MatrixFree`
- `MatrixFreeOperators` Idea:
  - Only overload write (cell) operator (`apply_add`),
  - `MatrixFreeOperators` takes care of the rest

Similar to `MeshWorker`.

All of them only use CPUs

## SPMV matrix-based vs. matrix-free

matrix-based

$$A = \sum_k P_k^T A_k P_k$$

$$v = Au$$

Separate steps for

- Assembly
- matrix-vector product

bandwidth limited: <1 Flop/Byte

matrix-free

$$v = A(u) = \sum_k^{ncells} P_k^T A_k P_k u$$

- Extract local elements:  
 $u_k = P_k u$
- Compute the matrix-vector product on the fly:  $v_k = A_k u_k$
- Add to global vector:  
 $v += P_k^T v_k$

## Matrix-vector product on cell, Laplacian

$$\begin{aligned}(A_K u_K)_j &= \int_K \nabla_x \phi \cdot \nabla_x u^h dx \approx \sum_q w_q \det J_q \nabla \phi_j \cdot \nabla u^h \Big|_{x=x_q} \\ &= \sum_q \nabla_\eta \phi_j J_q^{-1} (w_q \det J_q) J_q^{-T} \sum_i \nabla_\eta \phi_i u_{K,i} \Big|_{x=x_q}\end{aligned}$$

- Compute unit cell gradients  $\sum_i \nabla_\eta \phi_i u_{K,i}$  at all quadrature points
- At each quadrature point, apply geometry  $J_q^{-T}$ , multiply by quadrature weight and Jacobian determinant  $w_q \det J_q$ , and apply geometry for test function  $J_q^{-1}$
- Test by unit cell gradients of all basis functions and sum over quadrature points

## Sum factorization

$$\sum_i \phi_i(x_q) u_{K,i} = \sum_{i_z} \phi_{i_z}(x_{q_z}) \sum_{i_y} \phi_{i_y}(x_{q_y}) \sum_{i_x} \phi_{i_x}(x_{q_x}) u_{K,i}$$

Take advantage of the tensor product structure of shape functions and quadrature points since  $\xi_q = \xi_{q_x}, \xi_{q_y}, \xi_{q_z}$  and  $\phi_{x,y,z} = \phi_x \phi_y \phi_z$

$$S = \nabla_{\eta} \phi_i|_{\xi_q} = \begin{bmatrix} D_z \otimes D_y \otimes S_x \\ D_z \otimes S_y \otimes D_x \\ S_z \otimes D_y \otimes D_x \end{bmatrix}$$

with

$$D_x = \phi_x(\xi_{q_x})$$

$$S_x = \nabla_{\eta} \phi_x(\xi_{q_x})$$

Evaluation costs reduce from  $\mathcal{O}((p+1)^{2d})$  to  $\mathcal{O}(d(p+1)^{d+1})$ .

## History of Portable MatrixFree in deal.II

- introduced to deal.II with  
<https://github.com/dealii/dealii/pull/4293> (deal.II 9.0.0)
- decided to switch to Kokkos in  
<https://github.com/dealii/dealii/pull/15200> last year
- renamed in <https://github.com/dealii/dealii/pull/16497> this year

# Portable MatrixFree limitations

- Only FE\_Q elements, continuous Galerkin
- no face terms
- only one component
- limited optimizations, even-odd, cartesian meshes
- degree coupled with # quadrature points

# Portable MatrixFree implementation

## Kokkos implementation

- Uses TeamPolicy for looping over the cells (coloring possible), every team gets a cell batch
- scratch memory for working values and gradients buffers
- TeamThreadRange for tensor contractions on every cell, parallelize over quadrature points
- TeamThreadMDRange for newer Kokkos

## deal.II tutorial: step-64

In this example, we consider the Helmholtz problem

$$\begin{aligned} -\nabla \cdot \nabla u + a(\mathbf{x})u &= 1, \\ u &= 0 \quad \text{on } \partial\Omega \end{aligned}$$

where  $a(\mathbf{x})$  is a variable coefficient.

We choose as domain  $\Omega = [0, 1]^3$  and  $a(\mathbf{x}) = \frac{10}{0.05 + 2\|\mathbf{x}\|^2}$ . Since the coefficient is symmetric around the origin but the domain is not, we will end up with a non-symmetric solution.



## step-64

```
template <int dim, int fe_degree>
void HelmholtzOperator<dim, fe_degree>::vmult(
    LinearAlgebra::distributed::Vector<double, MemorySpace::Default> &dst,
    const LinearAlgebra::distributed::Vector<double, MemorySpace::Default> &src)
const
{
    dst = 0.;
    LocalHelmholtzOperator<dim, fe_degree> helmholtz_operator(
        coef.get_values());
    mf_data.cell_loop(helmholtz_operator, src, dst);
    mf_data.copy_constrained_values(src, dst);
}
```

## step-64

```
template <int dim, int fe_degree>
DEAL_II_HOST_DEVICE void LocalHelmholtzOperator<dim, fe_degree>::operator()(
    const unsigned int cell,
    const typename Portable::MatrixFree<dim, double>::Data *gpu_data,
    Portable::SharedData<dim, double> *shared_data,
    const double *src,
    double *dst) const
{
    Portable::FEEvaluation<dim, fe_degree, fe_degree + 1, 1, double> fe_eval(
        gpu_data, shared_data);
    fe_eval.read_dof_values(src);
    fe_eval.evaluate(EvaluationFlags::values | EvaluationFlags::gradients);
    fe_eval.apply_for_each_quad_point(
        HelmholtzOperatorQuad<dim, fe_degree>(gpu_data, coef, cell));
    fe_eval.integrate(EvaluationFlags::values | EvaluationFlags::gradients);
    fe_eval.distribute_local_to_global(dst);
}
```

## step-64

```
template <int dim, int fe_degree, int n_q_points_1d, int n_components_, typename Number>
DEAL_II_HOST_DEVICE void
FEEvaluation<dim, fe_degree, n_q_points_1d, n_components_, Number>::
    read_dof_values(const Number *src)
{
    Kokkos::parallel_for(
        Kokkos::TeamThreadRange(shared_data->team_member, n_q_points),
        [&](const int &i) {
            shared_data->values(i) = src[data->local_to_global(cell_id, i)];
        });
    shared_data->team_member.team_barrier();

    internal::resolve_hanging_nodes<dim, fe_degree, false>(
        shared_data->team_member,
        data->constraint_weights,
        data->constraint_mask(cell_id),
        shared_data->values);
}
```

## Generic - evaluate values and gradients in 3D

```
template <int direction, bool dof_to_quad, bool add, bool in_place,
          typename ViewTypeIn, typename ViewTypeOut>
DEAL_II_HOST_DEVICE void values(const ViewTypeIn in, ViewTypeOut out) const; //apply

values<0, true, false, true>(u, u);
team_member.team_barrier();
values<1, true, false, true>(u, u);
team_member.team_barrier();
values<2, true, false, true>(u, u);
team_member.team_barrier();

gradients<0, true, false, false>(u, Kokkos::subview(grad_u, Kokkos::ALL, 0));
gradients<1, true, false, false>(u, Kokkos::subview(grad_u, Kokkos::ALL, 1));
gradients<2, true, false, false>(u, Kokkos::subview(grad_u, Kokkos::ALL, 2));
```

## Generic

```
typename ViewTypeOut>
DEAL_II_HOST_DEVICE void apply_3d(
    const Kokkos::TeamPolicy<...>::member_type &team_member,
    const Kokkos::View<Number *, ...> shape_data, const ViewTypeIn in, ViewTypeOut out)
{ [...]
```

```
    Number t[n_q_points] = {};
    auto thread_policy = Kokkos::TeamThreadMDRange<Kokkos::Rank<3>, TeamType>(
        team_member, n_q_points_1d, n_q_points_1d, n_q_points_1d);
    Kokkos::parallel_for(
        thread_policy, [&](const int i, const int j, const int q) {
            const int q_point = ...;
            for (int k = 0; k < n_q_points_1d; ++k) {
                const unsigned int shape_idx = ...;
                const unsigned int source_idx = ...;
                t[q_point] += shape_data[shape_idx] * in(source_idx);
            }
        }); [...]
```

```
}
```

## step-64

```
template <int dim, int fe_degree>
DEAL_II_HOST_DEVICE void HelmholtzOperatorQuad<dim, fe_degree>::operator()(
    Portable::FEEvaluation<dim, fe_degree, fe_degree + 1, 1, double> *fe_eval,
    const int q_point) const
{
    const unsigned int pos =
        gpu_data->local_q_point_id(cell, n_q_points, q_point);

    fe_eval->submit_value(coef[pos] * fe_eval->get_value(q_point), q_point);
    fe_eval->submit_gradient(fe_eval->get_gradient(q_point), q_point);
}
```

## Generic integrate values and gradients in 3D

```
template <int direction, bool dof_to_quad, bool add, bool in_place,
          typename ViewTypeIn, typename ViewTypeOut>
DEAL_II_HOST_DEVICE void values(const ViewTypeIn in, ViewTypeOut out) const; // apply

gradients<2, false, true, false>(Kokkos::subview(grad_u, Kokkos::ALL, 2), u);
team_member.team_barrier();
gradients<1, false, true, false>(Kokkos::subview(grad_u, Kokkos::ALL, 1), u);
team_member.team_barrier();
gradients<0, false, true, false>(Kokkos::subview(grad_u, Kokkos::ALL, 0), u);
team_member.team_barrier();

values<2, false, false, true>(u, u);
team_member.team_barrier();
values<1, false, false, true>(u, u);
team_member.team_barrier();
values<0, false, false, true>(u, u);
team_member.team_barrier();
```

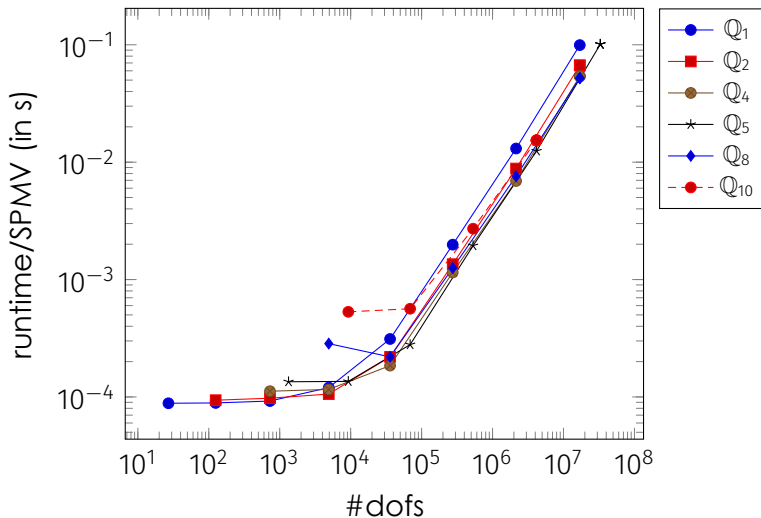
## step-64

```
template <int dim, int fe_degree, int n_q_points_1d, int n_components_, typename Number>
DEAL_II_HOST_DEVICE void
FEEvaluation<dim, fe_degree, n_q_points_1d, n_components_, Number>::
    distribute_local_to_global(Number *dst) const
{
    internal::resolve_hanging_nodes<dim, fe_degree, true>(
        shared_data->team_member, data->constraint_weights,
        data->constraint_mask(cell_id), shared_data->values);

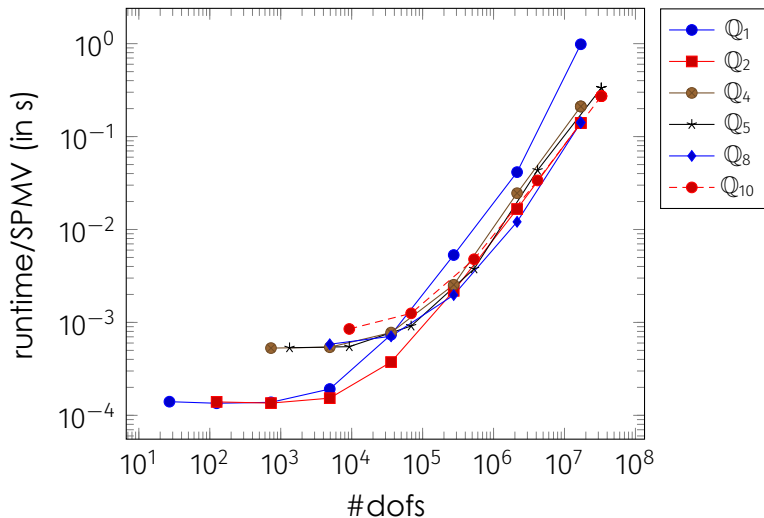
    if (data->use_coloring) {
        Kokkos::parallel_for(
            Kokkos::TeamThreadRange(shared_data->team_member, n_q_points),
            [&](const int &i) {
                dst[data->local_to_global(cell_id, i)] += shared_data->values(i);
            });
    }
    else { [...] }
}
```



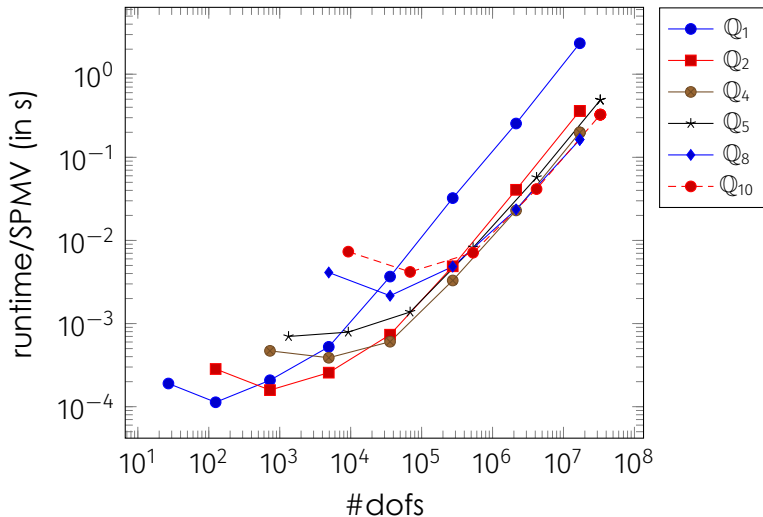
# Performance results on Polaris - NVIDIA A100



# Performance results on Frontier - AMD MI 250X, 1 GCD



# Performance results on Sunspot - Intel GPU Max 1550, 1 tile



# Summary

- GPUs can be used without changing code using Tpetra.
- Portable matrix-free implementation: we can run on CPUs, NVIDIA, AMD, and Intel GPUs using the same code
- Performance tuning needed, in particular for Intel GPUs
- Need to add more capabilities that the CPU version provides

# Questions?

# Acknowledgments

- This manuscript has been authored by UT-Battelle, LLC, under Contract No. DE-AC0500OR22725 with the U.S. Department of Energy.
- This work was funded by S4PST and OLCF through Kokkos Support funding.

## Cuda: Performance, A100

Results from bytes\_and\_flops(TeamPolicy)

scalar	Bandwidth	Compute	Cache
float	1251 GiB/s	14280 GFlop/s	3762 GiB/s
double	1267 GiB/s	7592 GFlop/s	6938 GiB/s
int32_t	1222 GiB/s	18457 GFlop/s	4684 GiB/s
int64_t	1267 GiB/s	3778 GFlop/s	6895 GiB/s

- Peak FP64 Vector: 19.5 TFLOPS
- Memory Bandwidth: 1.6 TB/sec
- Cache Size: 192KB (per SM)/40 MB

## HIP: Performance MI250, one GCD

Results from bytes\_and\_flops (TeamPolicy)

scalar	Bandwidth	Compute	Cache
float	1160 GiB/s	20544 GFlop/s	2756 GiB/s
double	1140 GiB/s	19320 GFlop/s	2883 GiB/s
int32_t	1150 GiB/s	20194 GFlop/s	2757 GiB/s
int64_t	1140 GiB/s	4979 GFlop/s	2865 GiB/s

- Peak FP64 Vector: 23.95 TFLOPS
- Memory Bandwidth: 1.6 TB/sec
- Cache Size: 16KB (per CU)/16 MB



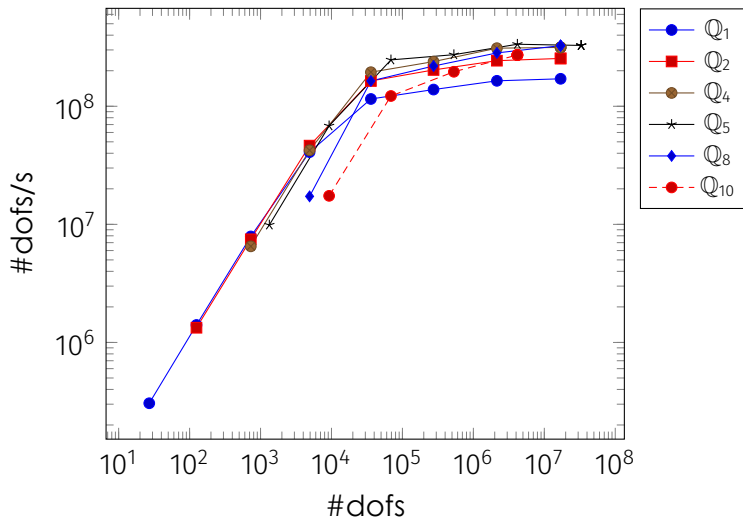
## SYCL: Performance Intel Data Center GPU Max 1550, one tile

Results from bytes\_and\_flops(TeamPolicy)

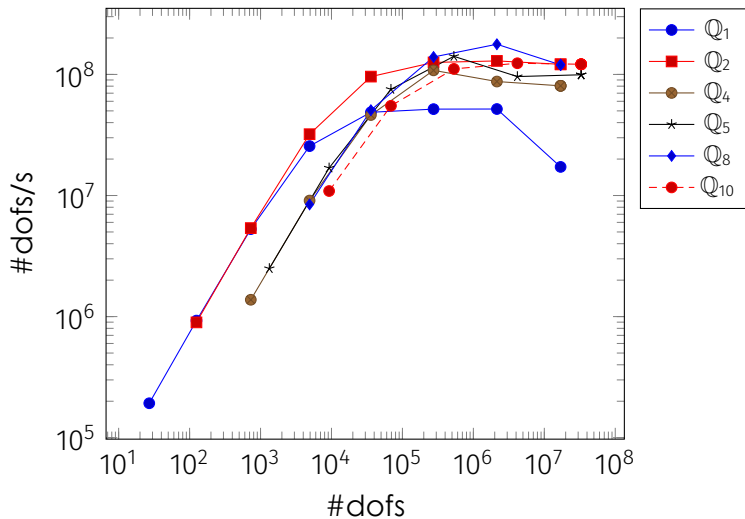
scalar	Bandwidth	Compute	Cache
float	1002 GiB/s	17484 GFlop/s	4973 GiB/s
double	960 GiB/s	8746 GFlop/s	6928 GiB/s
int32_t	1007 GiB/s	6108 GFlop/s	4714 GiB/s
int64_t	958 GiB/s	982 GFlop/s	4715 GiB/s

- Peak FP64 Vector: 22.9 TFLOPS/tile
- Memory Bandwidth: 1.6 TB/sec/tile
- Cache Size: 128KB (per work group)/408 MB

# Performance results on Polaris - NVIDIA A100



# Performance results on Frontier - AMD MI 250X, 1 GCD



# Performance results on Sunspot - Intel GPU Max 1550, 1 tile

